

IX-Randomness

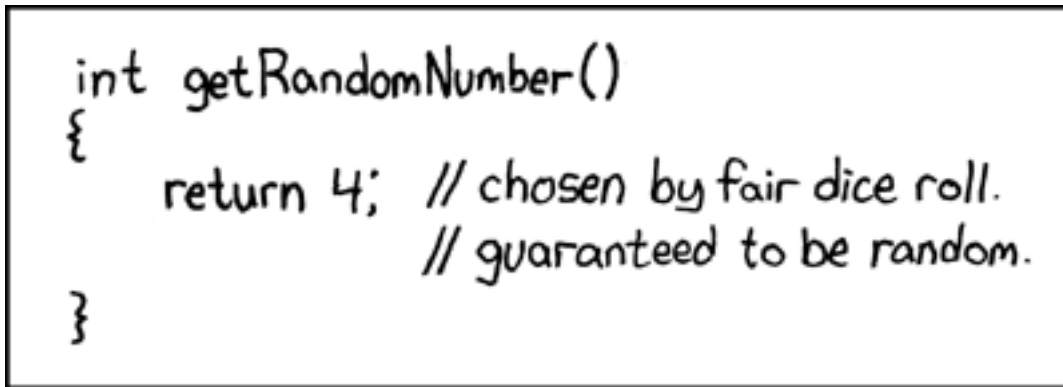
December 7, 2014

1 IX-Randomness

In [1]: `from IPython.display import Image`

In [2]: `Image('http://imgs.xkcd.com/comics/random_number.png')`

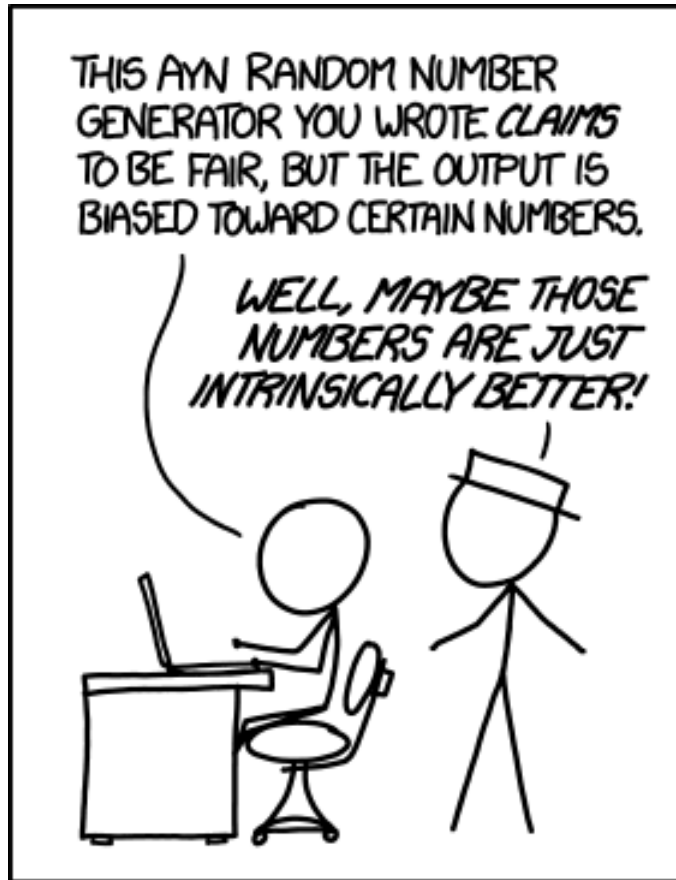
Out[2]:



```
int getRandomNumber()
{
    return 4; // chosen by fair dice roll.
              // guaranteed to be random.
}
```

In [3]: `Image('http://imgs.xkcd.com/comics/any_random.png')`

Out[3]:



```
In [4]: Image('http://www.random.org/analysis/dilbert.jpg')
```

Out[4]:



```
In [15]: import numpy as np
```

The *random* subpackage is the package responsible for all random functions in the numpy environment

```
In [16]: #help(np.random)
```

Let's generate a couple of random paths

```
In [26]: Nt = 1000
         Npaths = 10
         rv = np.random.randn(10,1000)
```

```
In [27]: rv.shape
```

```
Out[27]: (10, 1000)
```

```
In [28]: paths = rv.cumsum(1)
```

```
In [29]: paths.shape
```


```
Out[29]: (10, 1000)
```

```
In [21]: %matplotlib inline
         from matplotlib.pyplot import *
         import seaborn as sns
         sns.set(style='ticks', palette='Set2')
```

```
In [22]: plot(paths.T)
```

```
Out[22]: [<matplotlib.lines.Line2D at 0x7f09d8b351d0>,
         <matplotlib.lines.Line2D at 0x7f09d8b23650>,
         <matplotlib.lines.Line2D at 0x7f09d8b238d0>,
         <matplotlib.lines.Line2D at 0x7f09d8b23110>,
         <matplotlib.lines.Line2D at 0x7f09d8a9f110>,
         <matplotlib.lines.Line2D at 0x7f09d8a9f2d0>,
         <matplotlib.lines.Line2D at 0x7f09d8b3e6d0>,
         <matplotlib.lines.Line2D at 0x7f09d8a9f650>,
         <matplotlib.lines.Line2D at 0x7f09d8a9f810>,
         <matplotlib.lines.Line2D at 0x7f09d8a9f9d0>]
```

```
/home/jpsilva/anaconda/lib/python2.7/site-packages/matplotlib/font_manager.py:1279: UserWarning: findfor
(prop.get_family(), self.defaultFamily[fonttext]))
```



```
In [38]: from scipy import stats
```

We can easily get a list of available distributions (continuous and discrete) by using *dir* or by checking on the online [documentation](#)

```
In [45]: list_distributions = dir(stats)
        len(list_distributions)
```

```
Out[45]: 243
```

Let's choose a Pareto distribution with parameter $\alpha = 3$


$$pdf_{pareto}(x) = \frac{\alpha}{x^{\alpha+1}}$$

```
In [51]: alpha = 3
        X = stats.pareto(alpha)
```

We can sample...

```
In [50]: n = 100
        plot(X.rvs(n), '.')
```


```
Out[50]: [<matplotlib.lines.Line2D at 0x7f09d6ce03d0>]
```



...have a look at the probability density function...

```
In [58]: x = np.linspace(X.ppf(0.01), X.ppf(0.99), 100)
        plot(x, X.pdf(x), label='Pareto pdf')
        legend()
        xlabel('$x$')
```


Out[58]: <matplotlib.text.Text at 0x7f09d691bf90>



...have a look at the cumulative density function...

```
In [60]: plot(x, X.cdf(x),label='Pareto cdf')
         legend()
         xlabel('$x$')
```


Out[60]: <matplotlib.text.Text at 0x7f09d67ef390>



...compare the histogram of the previously generated sample with the pdf...

```
In [68]: fig, ax = subplots()
         ax.hist(X.rvs(10000),normed=True,bins=100)
         ax.plot(x,X.pdf(x))

Out[68]: [<matplotlib.lines.Line2D at 0x7f09d6016590>]
```



(let's just see the influence of the number of bins in the non-parametric estimation)

```
In [86]: %matplotlib
from matplotlib.widgets import Slider

samples = X.rvs(10000)

ax = axes([0.1,0.25,0.8,0.6])
ax.hist(samples,normed=True)
ax.plot(x,X.pdf(x))
sl = Slider(axes([0.1,0.1,0.8,0.1]),'Bins',1,200,valinit=10,valfmt='%d')

def update(data):
    data = int(data)
    ax.cla()
    ax.hist(samples,normed=True,bins=data)
    ax.plot(x,X.pdf(x))
    ax.draw()

sl.on_changed(update)
```

Using matplotlib backend: Qt4Agg

Out[86]: 0

```
In [87]: X.stats(moments='mv')
```

Out[87]: (array(1.5), array(0.75))

```
In [23]: from scipy.stats import kstest
```

```
In [30]: kstest(rv,'norm')
```

```
-----  
ValueError
```

```
Traceback (most recent call last)
```

```
<ipython-input-30-f62508e716a2> in <module>()  
----> 1 kstest(rv,'norm')
```

```
/home/jpsilva/anaconda/lib/python2.7/site-packages/scipy/stats/stats.pyc in kstest(rvs, cdf, ar  
3433  
3434     if alternative in ['two-sided', 'greater']:  
-> 3435         Dplus = (np.arange(1.0, N+1)/N - cdfvals).max()  
3436         if alternative == 'greater':  
3437             return Dplus, distributions.ksone.sf(Dplus,N)
```

```
ValueError: operands could not be broadcast together with shapes (10,) (10,1000)
```

```
In [31]: kstest(rv,'norm')
```

```
-----  
ValueError
```

```
Traceback (most recent call last)
```

```
<ipython-input-31-f62508e716a2> in <module>()  
----> 1 kstest(rv,'norm')
```

```
/home/jpsilva/anaconda/lib/python2.7/site-packages/scipy/stats/stats.pyc in kstest(rvs, cdf, ar  
3433  
3434     if alternative in ['two-sided', 'greater']:  
-> 3435         Dplus = (np.arange(1.0, N+1)/N - cdfvals).max()  
3436         if alternative == 'greater':  
3437             return Dplus, distributions.ksone.sf(Dplus,N)
```

```
ValueError: operands could not be broadcast together with shapes (10,) (10,1000)
```

```
In [33]: rv.apply(kstest,axis=1)
```

```
-----  
AttributeError
```

```
Traceback (most recent call last)
```

```
<ipython-input-33-0e6a60028561> in <module>()  
----> 1 rv.apply(kstest,axis=1)
```

```
AttributeError: 'numpy.ndarray' object has no attribute 'apply'
```



```

In [34]: np.apply_along_axis(lambda x: kstest(x, 'norm'), 1, rv)

Out[34]: array([[ 0.02462349,  0.57921451],
 [ 0.02313629,  0.65813139],
 [ 0.02690737,  0.46175665],
 [ 0.01927064,  0.85160859],
 [ 0.02455753,  0.58267687],
 [ 0.0551436 ,  0.00439168],
 [ 0.02582343,  0.51737188],
 [ 0.02793807,  0.41210544],
 [ 0.02163583,  0.73737606],
 [ 0.02397197,  0.61361516]])

In [35]: map(lambda x: kstest(x, 'norm'), rv)

Out[35]: [(0.024623490897107081, 0.57921450514349404),
 (0.02313628625770614, 0.65813139144955546),
 (0.026907370554564158, 0.46175664577782749),
 (0.019270644257675884, 0.85160859135592293),
 (0.02455753392874805, 0.58267687044366556),
 (0.055143604679474989, 0.0043916809490187614),
 (0.025823431496813432, 0.51737188246826193),
 (0.027938071951015719, 0.41210543959431489),
 (0.021635830354071017, 0.73737606358434737),
 (0.023971974347128389, 0.61361516398657845)]

In [36]: %%timeit
          np.apply_along_axis(lambda x: kstest(x, 'norm'), 1, rv)

100 loops, best of 3: 6.61 ms per loop

In [37]: %%timeit
          map(lambda x: kstest(x, 'norm'), rv)

100 loops, best of 3: 6.4 ms per loop

In [23]: np.apply_along_axis(??)

In [93]: t_statistic, p_value = stats.ttest_ind(X.rvs(size=1000), X.rvs(size=1000))
          print p_value

0.433472011446

In [90]:

In []:

In []:

In []:

In []:

In []:

In [25]: from os import urandom

In [26]: ur = urandom(16)

```

```
In [27]: ur.encode('hex')
```

```
Out[27]: 'a745c7bc678df931b1ac1945baf2c7d1'
```

```
In [28]: import random
```

```
In [29]: random.random()
```

```
Out[29]: 0.9962189820806353
```

What about choosing from some predetermined set?

```
In [1]: from os import urandom
```

```
In [9]: urandom(2).encode('hex')
```

```
Out[9]: '27b9'
```

```
In [10]: a = urandom(64)
         a.encode('base-64')
```

```
Out[10]: 'lLn50JCxw1DmD87MiiHHZDv0xjNpcG/h0uiGS2Y4kG6h1g0QaoySBKYiGX31SxD1weDkaFOYkTwa\nu1Akx92EHg==\n'
```

```
In [11]: from base64 import b64encode
         from os import urandom
```

```
         random_bytes = urandom(64)
         token = b64encode(random_bytes).decode('utf-8')
         token
```

```
Out[11]: u'H+y/fHdAumBMVJr3CKG6PJX/dg1su8WEc48z9I9MYECcOKDqiXz9+WoPS0/r1t/cJdyYzAvw68mPa+uFyEbpZw=='
```

```
In [3]: %load_ext version_information
         %version_information numpy
```

The version_information extension is already loaded. To reload it, use:
%reload_ext version_information

```
Out[3]:
```

Software	Version
Python	2.7.8 —Anaconda 2.1.0 (64-bit)— (default, Aug 21 2014, 18:22:21) [GCC 4.4.7 20120313 (Red Hat 4.4.7-1)]
IPython	2.3.0
OS	posix [linux2]
numpy	1.9.1
Thu Dec 04 14:48:14 2014 CET	

```
In [1]: from IPython.core.display import HTML
         def css_styling():
             styles = open("./styles/custom.css", "r").read()
             return HTML(styles)
         css_styling()
```

```
Out[1]: <IPython.core.display.HTML at 0x7fc5082118d0>
```

```
In []:
```